


Apache Flink waterMark的机制

• 本课目标

- ☑ 掌握WaterMark的原理
- ☑ 掌握WaterMark的运用

• 1. 案例需求

- 需求描述：每隔5秒，计算最近10秒单词出现的次数



- 1.1 TimeWindow实现

```
//案例1 (job1) -> timeWindow

public static void main(String[] args) throws Exception {
    StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();

    DataStreamSource<String> dataStream =
    env.socketTextStream("192.168.134.130", 8888);
    dataStream.flatMap(new FlatMapFunction<String, Tuple2<String,Integer>>()
    {
        @Override
        public void flatMap(String line, Collector<Tuple2<String, Integer>>
        out) throws Exception {
            String[] fields = line.split(",");
            for(String word:fields){
                out.collect(Tuple2.of(word,1));
            }
        }
    })
```

```

    }
    }).keyBy(0).timeWindow(Time.seconds(10),Time.seconds(5))
        .sum(1)
        .print().setParallelism(1);

    env.execute("TimeWindowWordCount");
}

```

- 1.2 TimeWindowProcessTest

```

//案例2 (job2) -> TimeWindowProcessTest
public static void main(String[] args) throws Exception {
    StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();

    DataStreamSource<String> dataStream =
        env.socketTextStream("192.168.134.130", 9999);
    dataStream.flatMap(new FlatMapFunction<String, Tuple2<String,Integer>>()
    {
        @Override
        public void flatMap(String line,
                               Collector<Tuple2<String, Integer>> out) throws
Exception {
            String[] fields = line.split(",");
            for(String word:fields){
                out.collect(Tuple2.of(word,1));
            }
        }
    }).keyBy(0).timeWindow(Time.seconds(10),Time.seconds(5))
        .process(new MySumProcessWindowFunction()) // 相当于spark里面的
foreach
        .print().setParallelism(1);

    env.execute("TimeWindowWordCount");
}

/**
 * <IN, OUT, KEY, W extends Window>
 *     IN: 输入的数据类型
 *     OUT: 输出的数据类型
 *     KEY: key的数据类型
 *     W: 窗口的数据类型
 *
 */
public static class MySumProcessWindowFunction extends
ProcessWindowFunction<Tuple2<String,Integer>,Tuple2<String,Integer>,Tuple,TimeWin
dow>{

```

```

FastDateFormat dataformat = FastDateFormat.getInstance("HH:mm:ss");

//elements: 当前窗口里面的数据
@Override
public void process(Tuple key, Context context,
                    Iterable<Tuple2<String, Integer>> elements,
                    Collector<Tuple2<String, Integer>> out) throws
Exception {

    System.out.println("当前系统时
间: "+dataformat.format(System.currentTimeMillis()));
    System.out.println("窗口处理时
间: "+dataformat.format(context.currentProcessingTime()));
    System.out.println("窗口开始时
间: "+dataformat.format(context.window().getStart()));

    int sum=0;
    for (Tuple2<String,Integer> ele:elements){
        sum += 1;
    }

    out.collect(Tuple2.of(key.getField(0),sum));
    System.out.println("窗口结束时
间: "+dataformat.format(context.window().getEnd()));

    System.out.println("=====");
}
}

```

运行结果:

```

当前系统时间: 09:57:35
窗口处理时间: 09:57:35
窗口开始时间: 09:57:25
窗口结束时间: 09:57:35
(hive,1)

```

```

当前系统时间: 09:57:40
窗口处理时间: 09:57:40
窗口开始时间: 09:57:30
窗口结束时间: 09:57:40
(hive,2)

```

```

当前系统时间: 09:57:45
窗口处理时间: 09:57:45
窗口开始时间: 09:57:35
窗口结束时间: 09:57:45
(hive,1)

```

```

当前系统时间: 09:57:50
窗口处理时间: 09:57:50

```

窗口开始时间: 09:57:40

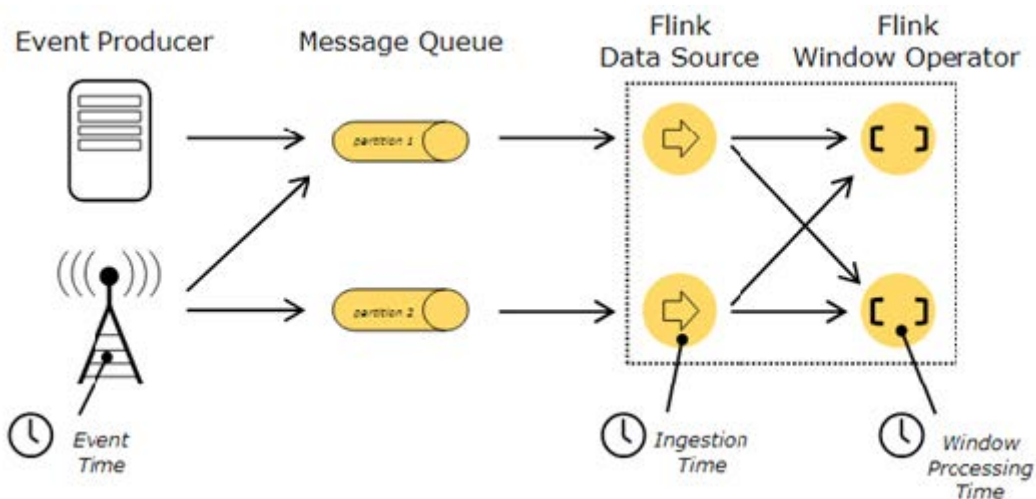
窗口结束时间: 09:57:50

根据每隔5秒执行最近10秒的数据, Flink划分的窗口

[00:00:00, 00:00:05) [00:00:05, 00:00:10)
[00:00:10, 00:00:15) [00:00:15, 00:00:20)
[00:00:20, 00:00:25) [00:00:25, 00:00:30)
[00:00:30, 00:00:35) [00:00:35, 00:00:40)
[00:00:40, 00:00:45) [00:00:45, 00:00:50)
[00:00:50, 00:00:55) [00:00:55, 00:01:00)
[00:01:00, 00:01:05) ...

• 2. Time的种类

- 针对stream数据中的时间, 可以分为以下三种:
 - Event Time: 事件产生的时间, 它通常由事件中的时间戳描述
 - Ingestion time: 事件进入Flink的时间
 - Processing Time: 事件被处理时当前系统的时间 (默认)



案例演示: 原始日志如下

2021-10-10 13:00:01,134 INFO executor.Executor: Finished task in state 0.0

2021-10-10 13:00:01,134 是Event time

2021-10-10 20:00:03,102 是Ingestion time


2021-10-10 20:00:05,100 是Processing time

思考:

如果我们想要统计每分钟接口调用失败的错误日志个数, 使用哪个时间才有意义?

• 3.Process Time Window (有序)

- 需求: 每隔5秒计算最近10秒的单词出现的次数
- 自定义source, 模拟: 第 13 秒的时候连续发送 2 个事件, 第 16 秒的时候再发送 1 个事件



```
//案例3 -> job3
public static void main(String[] args) throws Exception {
    StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();

    DataStreamSource<String> dataStream = env.addSource(new TestSource());
    dataStream.flatMap(new FlatMapFunction<String, Tuple2<String,Integer>>()
{
    @Override
    public void flatMap(String line, Collector<Tuple2<String, Integer>>
out) throws Exception {
        String[] fields = line.split(",");
        for (String word:fields){
            out.collect(Tuple2.of(word,1));
        }
    }
}))

    .keyBy(0)
    .timeWindow(Time.seconds(10),Time.seconds(5))
    .process(new SumProcessFunction()).print().setParallelism(1);
    env.execute("WindowWordCountSortTest");
}

public static class TestSource implements SourceFunction<String> {
    FastDateFormat dateformat = FastDateFormat.getInstance("HH:mm:ss");

    @Override
    public void run(SourceContext<String> cxt) throws Exception {

        String currTime = String.valueOf(System.currentTimeMillis());

        System.out.println(currTime);

        //这个操作是我为了保证是10seconds的倍数。
        while(Integer.valueOf(currTime.substring(currTime.length() - 4)) >
10){

            currTime=String.valueOf(System.currentTimeMillis());
            continue;
        }
    }
}
```

```

        System.out.println("开始发送事件的时间: "+dateformat.format(System.currentTimeMillis()));
        TimeUnit.SECONDS.sleep(3);
        cxt.collect("flink");
        cxt.collect("flink");

        TimeUnit.SECONDS.sleep(3);
        cxt.collect("flink");

        //程序等待
        TimeUnit.SECONDS.sleep(500000);
    }

    @Override
    public void cancel() {

    }
}

/**
 * IN
 * OUT
 * KEY -> is Tuple
 * W extends Window
 *
 */
public static class SumProcessFunction extends
ProcessWindowFunction<Tuple2<String,Integer>,Tuple2<String,Integer>,Tuple,TimeWin
dow>{

    FastDateFormat dataformat=FastDateFormat.getInstance("HH:mm:ss");

    @Override
    public void process(Tuple tuple, Context context,
                       Iterable<Tuple2<String, Integer>> allElements,
                       Collector<Tuple2<String, Integer>> out) {
        int count = 0;
        for (Tuple2<String,Integer> e:allElements){
            count++;
        }
        out.collect(Tuple2.of(tuple.getField(0),count));
    }
}


```

输出结果:

```

(flink 2)
(flink 3)
(flink 1)

```



• 4. Process Time Window (无序)

- 自定义source模拟：第 13 秒的时候连续发送2个事件，第一个事件在第13秒的时候发送出去了，第二个事件在13秒的时候产生因为网络延迟等原因，在19秒的时候才发送出去，第三个事件16秒的时候发送了出去。

```
//案例4 -> job4
public static void main(String[] args) throws Exception {
    StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();

    DataStreamSource<String> dataStream = env.addSource(new TestSource());
    dataStream.flatMap(new FlatMapFunction<String, Tuple2<String,Integer>>()
    {
        @Override
        public void flatMap(String line,
                             Collector<Tuple2<String, Integer>> out) throws
    Exception {
        String[] fields = line.split(",");
        for (String word:fields){
            out.collect(Tuple2.of(word,1));
        }
    }
    })

    .keyBy(0)
    .timeWindow(Time.seconds(10),Time.seconds(5))
    .process(new SumProcessFunction()).print().setParallelism(1);

    env.execute("WindowBySort2");
}

public static class TestSource implements
    SourceFunction<String>{
    FastDateFormat dateformat = FastDateFormat.getInstance("HH:mm:ss");
    @Override
```

```

        public void run(SourceContext<String> cxt) throws Exception {
            String currTime = String.valueOf(System.currentTimeMillis());
            while(Integer.valueOf(currTime.substring(currTime.length() - 4)) >
100){
                currTime=String.valueOf(System.currentTimeMillis());
                continue;
            }
            System.out.println("开始发送事件的时间: "+dateformat.format(System.currentTimeMillis()));

            TimeUnit.SECONDS.sleep(3);
            //实际上我们的数据是在13秒的时候生成的，只是19的时候被发送出去。
            String event = "flink";
            //13s 第一个事件发送
            cxt.collect(event);

            //16s 第三个事件发送
            TimeUnit.SECONDS.sleep(3);
            cxt.collect("flink");

            //19s 第三个事件发送
            TimeUnit.SECONDS.sleep(3);
            cxt.collect(event);
            TimeUnit.SECONDS.sleep(500000);
        }

        @Override
        public void cancel() {

        }

    }

    /**
     * IN
     * OUT
     * KEY
     * W extends Window
     *
     */
    public static class SumProcessFunction
        extends
ProcessWindowFunction<Tuple2<String,Integer>, Tuple2<String,Integer>, Tuple, TimeWin
dow>{

        FastDateFormat dataformat=FastDateFormat.getInstance("HH:mm:ss");
        @Override
        public void process(Tuple tuple, Context context,
                            Iterable<Tuple2<String, Integer>> allElements,
                            Collector<Tuple2<String, Integer>> out) {

            int count=0;
            for (Tuple2<String,Integer> e:allElements){
                count++;
            }
        }
    }

```



```

    }
    out.collect(Tuple2.of(tuple.getField(0),count));
  }
}

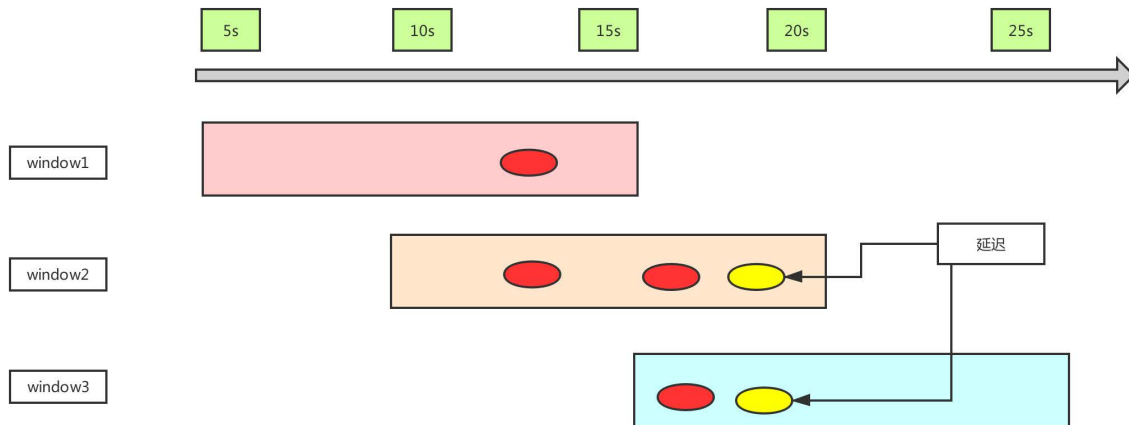
```

结果:

```

(flink,1)
(flink,3)
(flink,2)

```



• 5. 使用Event Time处理无序

//案例5 -> job5

```

public static void main(String[] args) throws Exception {
    StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();

    //步骤一: 设置时间类型, 默认的是Processtime
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStreamSource<String> dataStream = env.addSource(new TestSource());
    dataStream.map(new MapFunction<String, Tuple2<String,Long>>() {
        @Override
        public Tuple2<String, Long> map(String line) throws Exception {
            String[] fields = line.split(",");
            //key/value
            return new Tuple2<>(fields[0],Long.valueOf(fields[1]));
        }
    }).assignTimestampsAndWatermarks(new EventTimeExtractor())
        .keyBy(0)
        .timeWindow(Time.seconds(10),Time.seconds(5))
        .process(new SumProcessFunction()).print().setParallelism(1);

    env.execute("WindowWordCountAndTime");
}

```

```

    }

    public static class TestSource implements SourceFunction<String>{
        FastDateFormat dateformat = FastDateFormat.getInstance("HH:mm:ss");
        @Override
        public void run(SourceContext<String> cxt) throws Exception {
            String currTime = String.valueOf(System.currentTimeMillis());
            while(Integer.valueOf(currTime.substring(currTime.length() - 4)) >
100){
                currTime=String.valueOf(System.currentTimeMillis());
                continue;
            }
            System.out.println("开始发送事件的时间: "+dateformat.format(System.currentTimeMillis()));
            TimeUnit.SECONDS.sleep(3);
            //13s 第1个事件
            String event="flink,"+System.currentTimeMillis();//时间
            cxt.collect(event);

            TimeUnit.SECONDS.sleep(3);//16s 第2个事件
            cxt.collect("flink,"+System.currentTimeMillis());

            TimeUnit.SECONDS.sleep(3);
            //19s 第3个事件
            cxt.collect(event);

            TimeUnit.SECONDS.sleep(3000);

        }

        @Override
        public void cancel() {

        }
    }

    private static class EventTimeExtractor
        implements AssignerWithPeriodicWatermarks<Tuple2<String,Long>>{
        @Nullable
        @Override
        public Watermark getCurrentWatermark() {
            return new Watermark(System.currentTimeMillis());
        }

        //指定时间
        @Override
        public long extractTimestamp(Tuple2<String, Long> element, long l) {

```

```

        return element.f1;
    }
}

/**
 * IN
 * OUT
 * KEY
 * W extends Window
 */
public static class SumProcessFunction
    extends
ProcessWindowFunction<Tuple2<String,Long>,Tuple2<String,Integer>,Tuple,TimeWindow
>{

    FastDateFormat dataformat=FastDateFormat.getInstance("HH:mm:ss");
    @Override
    public void process(Tuple tuple, Context context,
                        Iterable<Tuple2<String, Long>> allElements,
                        Collector<Tuple2<String, Integer>> out) {
        int count=0;
        for (Tuple2<String,Long> e:allElements){
            count++;
        }
        out.collect(Tuple2.of(tuple.getField(0),count));
    }
}
}

```

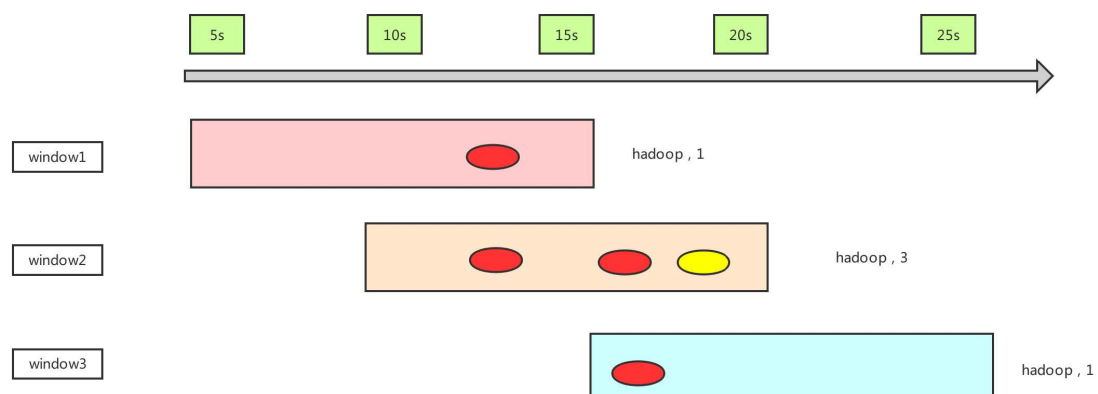
运行结果：但是存在数据丢失的问题

(flink,1)

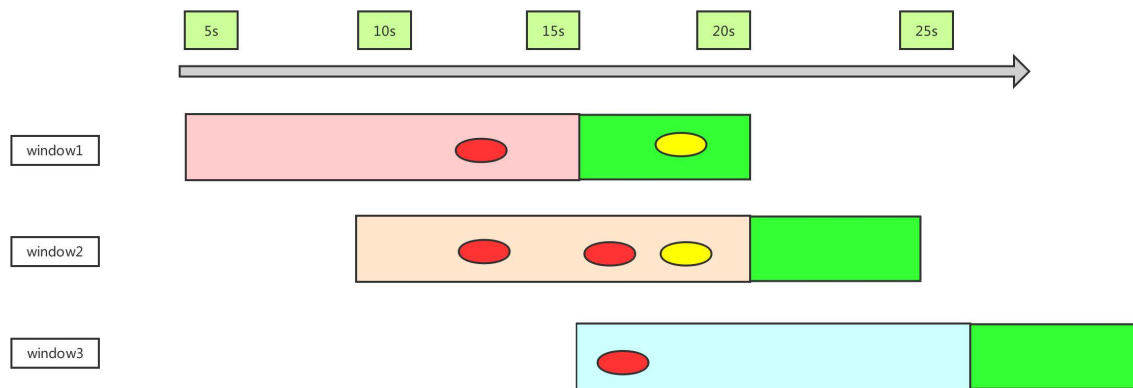
(flink,3)

(flink,1)

PS：现在我们第三个window的结果已经计算准确了，但是我们还是没有彻底的解决问题。接下来就需要我们使用WaterMark机制来解决了。



• 6. 使用WaterMark机制解决无序



//案例6 -> job6

```
public static void main(String[] args) throws Exception {
    StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();

    //步骤一：设置时间类型，默认的是Processtime
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStreamSource<String> dataStream = env.addSource(new TestSource());
    dataStream.map(new MapFunction<String, Tuple2<String,Long>>() {
        @Override
        public Tuple2<String, Long> map(String line) throws Exception {
            String[] fields = line.split(",");
            return new Tuple2<>(fields[0],Long.valueOf(fields[1]));
        }
    }).assignTimestampsAndWatermarks(new EventTimeExtractor())
        .keyBy(0)
        .timeWindow(Time.seconds(10),Time.seconds(5))
        .process(new SumProcessFunction()).print().setParallelism(1);

    env.execute("WindowWordCountAndTime");
}

public static class TestSource implements SourceFunction<String>{
    FastDateFormat dateformat = FastDateFormat.getInstance("HH:mm:ss");
    @Override
    public void run(SourceContext<String> cxt) throws Exception {
        String currTime = String.valueOf(System.currentTimeMillis());
        while(Integer.valueOf(currTime.substring(currTime.length() - 4)) >
100){
            currTime=String.valueOf(System.currentTimeMillis());
            continue;
        }
        System.out.println("开始发送事件的时间: "+dateformat.format(System.currentTimeMillis()));
        TimeUnit.SECONDS.sleep(3);
    }
}
```

```

        String event="flink,"+System.currentTimeMillis();
        cxt.collect(event);

        TimeUnit.SECONDS.sleep(3);
        cxt.collect("flink,"+System.currentTimeMillis());

        TimeUnit.SECONDS.sleep(3);
        cxt.collect(event);
        TimeUnit.SECONDS.sleep(500000);
    }

    @Override
    public void cancel() {

    }
}

private static class EventTimeExtractor implements
AssignerWithPeriodicWatermarks<Tuple2<String,Long>>{
    //设置 5s的延迟 (乱序)
    @Nullable
    @Override
    public Watermark getCurrentWatermark() {
        //          System.out.println("water
maker:....."+System.currentTimeMillis());
        return new Watermark(System.currentTimeMillis() - 5000);
    }

    @Override
    public long extractTimestamp(Tuple2<String, Long> element, long l) {
        return element.f1;
    }
}

/**
 * IN
 * OUT
 * KEY
 * W extends Window
 *
 */
public static class SumProcessFunction
    extends
ProcessWindowFunction<Tuple2<String,Long>,Tuple2<String,Integer>,Tuple,TimeWindow
>{

    FastDateFormat dataformat=FastDateFormat.getInstance("HH:mm:ss");
    @Override
    public void process(Tuple tuple, Context context,
                       Iterable<Tuple2<String, Long>> allElements,
                       Collector<Tuple2<String, Integer>> out) {

```

```

    int count=0;
    for (Tuple2<String,Long> e:allElements){
        count++;
    }
    out.collect(Tuple2.of(tuple.getField(0),count));
}
}

```

• 7. WaterMark机制

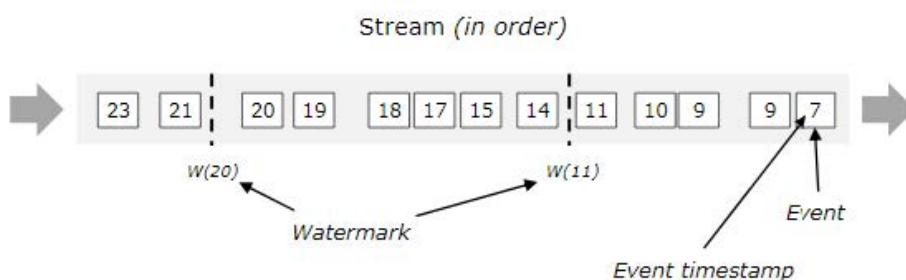
“

使用eventTime的时候如何处理乱序数据？我们知道，流处理从事件产生，到流经source，再到operator，中间是有一个过程和时间的。虽然大部分情况下，流到operator的数据都是按照事件产生的时间顺序来的，但是也不排除由于网络延迟等原因，导致乱序的产生，特别是使用kafka的话，多个分区的数据无法保证有序。所以在进行window计算的时候，我们又不能无限期的等下去，必须要有个机制来保证一个特定的时间后，必须触发window去进行计算了。这个特别的机制就watermark，watermark是用于处理乱序事件的。watermark可以翻译为水位线。

WaterMark：是周期性的运行的。在一个窗口周期当中

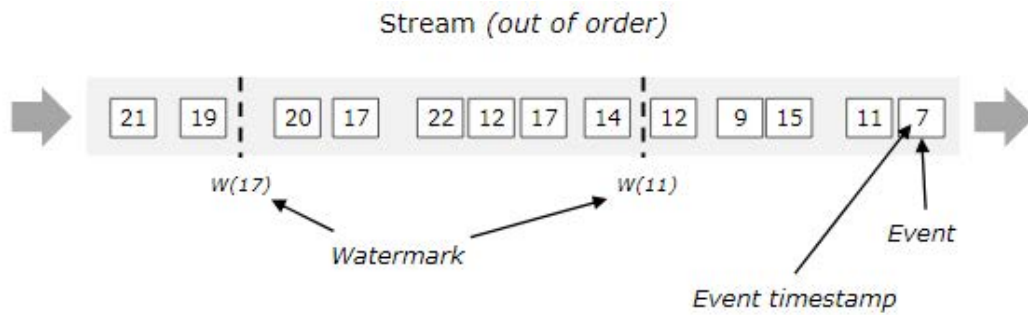
- 7.1 有序的流的watermarks

- 到达w(11)，一定会触发右边的window执行。
- 到达w(20)，一定会触发右边的window执行。
- 如果数据是有序的话，watermark意义不大




- 7.2 无序的流的watermarks

- w(11)：小于11的数据开始进行window计算，12，15不参与计算。
- w(17)：小于17的数据开始进行window计算



- 7.3 案例需求

- 得到并打印每隔 3 秒钟统计前 3 秒内的相同的 key 的所有的事件
- 9, 12, 14, 都比15小, 所以他们的watermark都是11, 因为当前事件的eventTime取的是区间的最大值.



案例7 -> job7

基于waterMark演示有序的数据计算

- 7.4 计算window的触发时间

- window触发的时间
 - watermark 时间 \geq window_end_time
 - 在 $[\text{window_start_time}, \text{window_end_time})$ 区间中有数据存在, 注意是左闭右开的区间, 而且是以 event time 来计算的
 - watermark出发时间 19:34:24, 这个时候刚好运算[19:34:21-19:34:24)产生的数据, 也就是运行10秒前产生的数据, 窗口为3秒。也就是根据当前的时间运行过去10秒之前的watermark。

Key	EventTime	CurrentMaxTimestamp	CurrentWaterMark	window_start_time	window_end_time
hadoop	19:34:22	19:34:22	19:34:12		
hadoop	19:34:26	19:34:26	19:34:16		
hadoop	19:34:32	19:34:32	19:34:22		
hadoop	19:34:33	19:34:33	19:34:23		
hadoop	19:34:34	19:34:34	19:34:24	[19:34:21	19:34:24)
hadoop	19:34:36	19:34:36	19:34:26		
hadoop	19:34:37	19:34:37	19:34:27	[19:34:24	19:34:27)

```

[00:00:00,00:00:03)
[00:00:03,00:00:06)
[00:00:06,00:00:09)
[00:00:09,00:00:12)
[00:00:12,00:00:15)
[00:00:15,00:00:18)
[00:00:18,00:00:21)
[00:00:21,00:00:24)
[00:00:24,00:00:27)
[00:00:27,00:00:30)
[00:00:30,00:00:33)
[00:00:33,00:00:36)
[00:00:36,00:00:39)
[00:00:39,00:00:42)
[00:00:42,00:00:45)
[00:00:45,00:00:48)
[00:00:48,00:00:51)
[00:00:51,00:00:54)
[00:00:54,00:00:57)
[00:00:57,00:01:00)
...

```

- 7.5 WaterMark+Window 处理乱序时间

案例7 -> job7

基于waterMark实现无序的数据计算


000001,1461756879000

000001,1461756871000

000001,1461756883000

• 7.6 事件迟到

- 丢弃，这个是默认的处理方式
- allowedLateness 指定允许数据延迟的时间
- sideOutputLateData 收集迟到的数据



- 7.6.1 丢弃

- 不推荐使用该方式

案例8 -> job8

重启程序，做测试。

输入数据：

000001,1461756870000

000001,1461756883000

000001,1461756870000

000001,1461756871000

000001,1461756872000

发现迟到太多数据就会被丢弃

- 7.6.2 再次推迟迟到的时间（了解）

- 当我们设置允许迟到 2 秒的事件，第一次 window 触发的条件是 watermark >= window_end_time
- 第二次(或者多次)触发的条件是 watermark < window_end_time + allowedLateness

案例9 -> job9

```
assignTimestampsAndWatermarks(new EventTimeExtractor())
    .keyBy(0)
    .timeWindow(Time.seconds(3))
    .allowedLateness(Time.seconds(2)) // 允许事件迟到 2 秒
```


- 7.6.3 收集迟到的数据（推荐）

案例9 -> job9

```
OutputTag<Tuple2<String, Long>> outputTag = new OutputTag<Tuple2<String, Long>>
("late-date"){};
.sideOutputLateData(outputTag) //保留迟到太多的数据
result.getSideOutput(outputTag).map(new MapFunction<Tuple2<String,Long>, String>
() {
    @Override
    public String map(Tuple2<String, Long> stringLongTuple2) throws
Exception {
        return "迟到数据: "+stringLongTuple2.toString();
    }
}).print();
```

• 8. 多并行度下的WaterMark

- 一个window可能会接受到多个waterMark，我们以最小的为准。



案例演示: job10

```
000001,1461756870000
000001,1461756883000
000001,1461756888000
```

当前线程ID: 55event = (000001,1461756883000)|19:34:43|19:34:43|19:34:33

当前线程ID: 56event = (000001,1461756870000)|19:34:30|19:34:30|19:34:20

当前线程ID: 56event = (000001,1461756888000)|19:34:48|19:34:48|19:34:38

处理时间: 19:31:25

window start time : 19:34:30

2> [(000001,1461756870000)|19:34:30]

window end time : 19:34:33

ID为56的线程有两个WaterMark: 20,38

那么38这个会替代20, 所以ID为56的线程的WaterMark是38

然后ID为55的线程的WaterMark是33, 而ID为56是WaterMark是38, 会在里面求一个小的值作为waterMark, 就是33, 这个时候会触发Window为30-33的窗口, 那这个窗口里面就有(000001,1461756870000)这条数据。